

Revisiting Rainbow: Promoting more insightful and inclusive deep reinforcement learning research



Johan S. Obando Ceron, Pablo Samuel Castro*
*Google Research, Brain Team




Overview:

Motivation: We argue for a need to change the status-quo in evaluating and proposing new research to avoid exacerbating the barriers to entry for newcomers from underprivileged communities.

- We complement this argument by revisiting the Rainbow algorithm on a set of small- and medium-sized tasks.
- This work allows to conduct a counterfactual analysis, and investigate whether there is scientific value in exploring empirical research in reinforcement learning when restricting oneself to small- to mid-scale environments.

Environments [1, 2]:



Results:

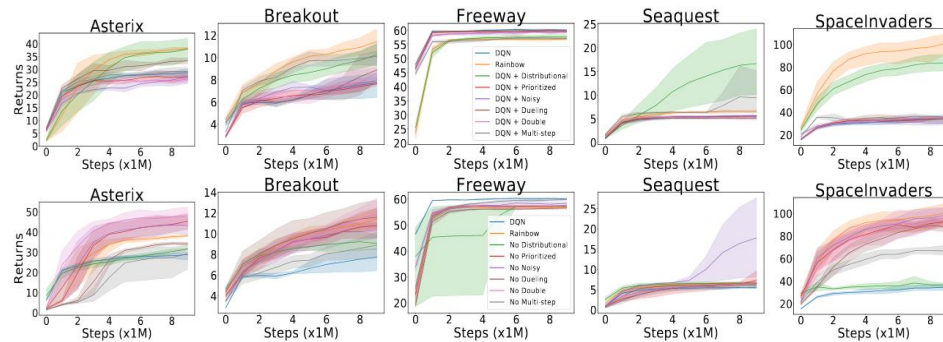


Fig.1. Comparison of the different algorithmic components on five MinAtar games, using the optimal hyper-parameters for each. Top row explores adding on top of DQN, bottom row explores removing from Rainbow.




Fig.2. Evaluating DQN sensitivity to varying number of layers (top) and units per layer (bottom).

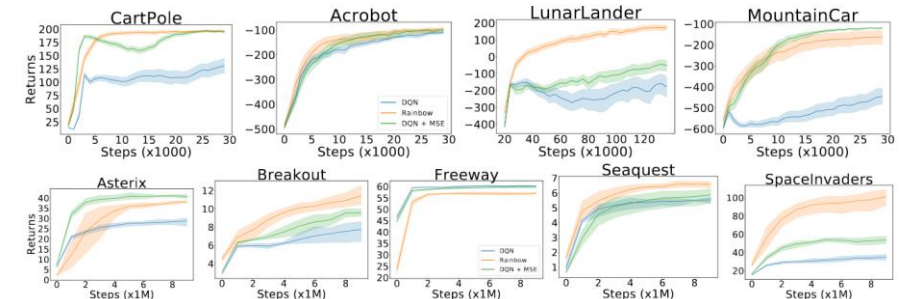


Fig.3. Evaluation of the use of the mean-squared error loss, instead of the Huber loss, in DQN.

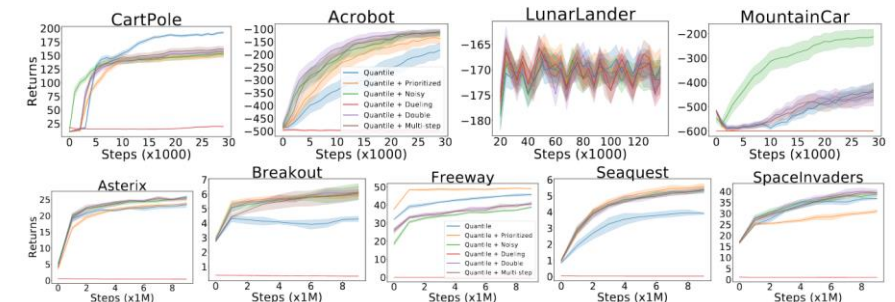


Fig.4. Investigating the combination of the different algorithmic components with QR-DQN.

References:

1. Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym, 2016.
2. Kenny Young and Tian Tian. Minatar: An atari-inspired testbed for thorough and reproducible reinforcement learning experiments. arXiv preprint arXiv:1903.03176, 2019.